

P8106_HW4_yz4184

Yunlin Zhou

Contents

Problem 1	2
Part a	2
Part b	5
Part c	8
Question 2	11
Part a	11
Part b	16

```
library(tidyverse)
library(ISLR)
library(caret)
library(rpart)
library(rpart.plot)
library(party)
library(partykit)
library(randomForest)
library(ranger)
library(gbm)
library(pdp)
library(pROC)
```

Problem 1

```
# Import and clean the data
college_df = read.csv("./College.csv")%>%
  janitor::clean_names()%>%
  drop_na()%>%
  relocate("outstate")%>%
  select(-college)

# Partition data into training/test sets
set.seed(1)
college_train = createDataPartition(y = college_df$outstate,
                                     p = 0.8,
                                     list = FALSE)

train_df = college_df[college_train,]
test_df = college_df[-college_train,]
```

Part a

Build a regression tree on the training data to predict the response.


```
ctrl <- trainControl(method = "cv")

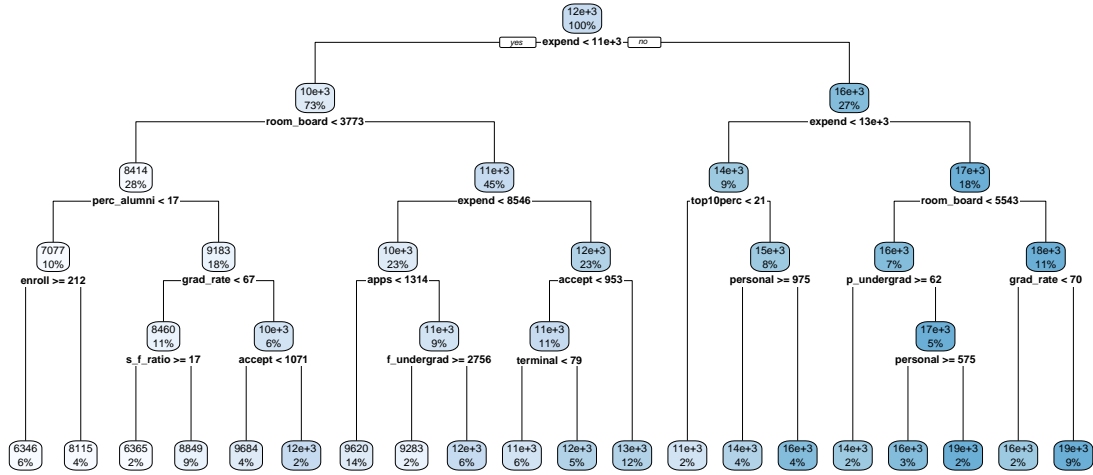
set.seed(1)
tree1 <- train(outstate ~ .,
               train_df,
               method = "rpart",
               tuneGrid = data.frame(cp = exp(seq(-6, -2, length = 50))),
               trControl = ctrl)

tree1$bestTune

##           cp
## 8 0.004389362
```

```
ggplot(tree1, highlight = TRUE)
```





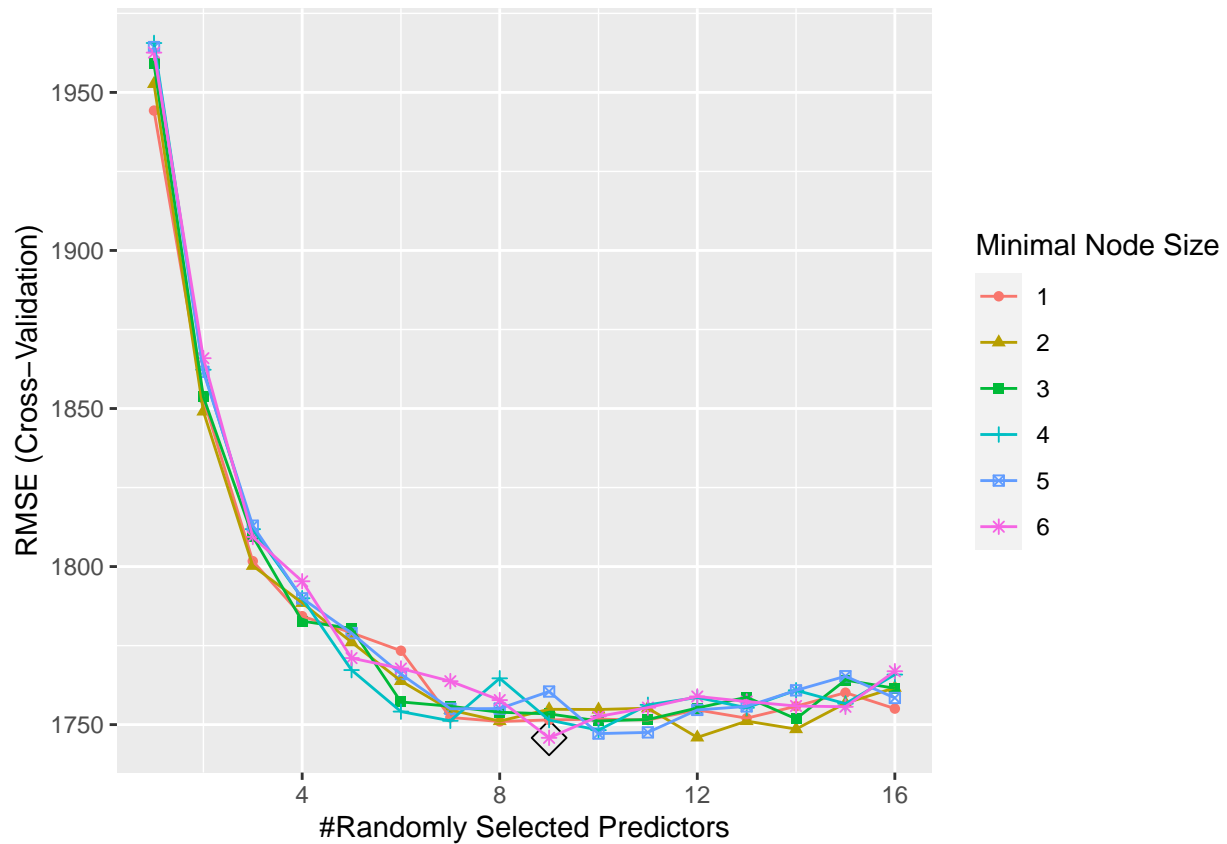
Part b

Perform random forest on the training data.

```
set.seed(1)
rf.grid <- expand.grid(mtry = 1:16,
                      splitrule = "variance",
                      min.node.size = 1:6)

rf1 <- train(outstate ~ .,
             train_df,
             method = "ranger",
             tuneGrid = rf.grid,
             trControl = ctrl)

ggplot(rf1, highlight = TRUE)
```



```
rf1$bestTune
```


```
##      mtry splitrule min.node.size
## 54      9  variance              6
```

In this random forest model, the best model is with minimum node size 6 and 9 selected predictors.

Report the variable importance.

```
set.seed(1)
rf1.final.per <- ranger(outstate ~ . ,
                        train_df,
                        mtry = rf1$bestTune[[1]],
                        splitrule = "variance",
                        min.node.size = rf1$bestTune[[3]],
                        importance = "permutation",
                        scale.permutation.importance = TRUE)

barplot(sort(ranger::importance(rf1.final.per), decreasing = FALSE),
        las = 2, horiz = TRUE, cex.names = 0.7,
        col = colorRampPalette(colors = c("cyan", "blue"))(16))
```



Using the permutation method, the most important predictors are `expend` and `room_board`.

Report the test error.

```
pred.rf <- predict(rf1, newdata = test_df)
te_rf = RMSE(pred.rf, test_df$outstate)
te_rf
```

```
## [1] 1651.307
```

The test error is 1651.3069135.

Part c

Perform boosting on the training data.

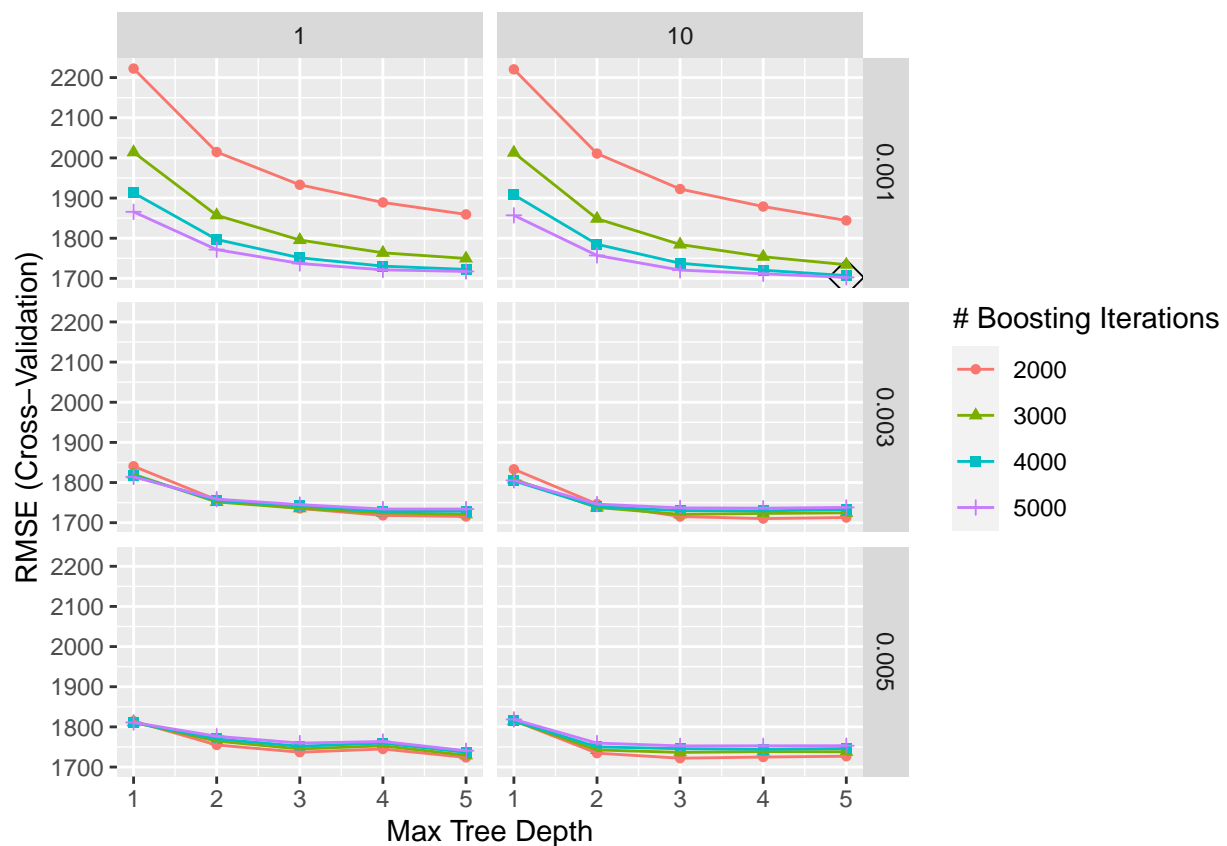
```
gbm.grid <- expand.grid(n.trees = c(2000,3000,4000,5000),
                      interaction.depth = 1:5,
                      shrinkage = c(0.001,0.003,0.005),
                      n.minobsinnode = c(1,10))

set.seed(1)
gbm1 <- train(outstate ~ . ,
             train_df,
             method = "gbm",
             tuneGrid = gbm.grid,
             trControl = ctrl,
             verbose = FALSE)

gbm1$bestTune
```


```
##      n.trees interaction.depth shrinkage n.minobsinnode
## 40      5000                5      0.001             10
```

```
ggplot(gbm1, highlight = TRUE)
```



Report the variable importance.

```
summary(gbm1$finalModel, las = 2, cBars = 16, cex.names = 0.6)
```



```
##           var  rel.inf
## expend      expend 56.839618
## room_board  room_board 11.763945
## grad_rate   grad_rate  4.097538
## terminal    terminal  3.740792
## perc_alumni perc_alumni 3.692448
## apps        apps     3.212812
## accept      accept    2.606760
## personal    personal  2.368107
## top10perc   top10perc  1.812568
## f_undergrad f_undergrad 1.605638
## ph_d        ph_d      1.562622
## s_f_ratio    s_f_ratio  1.445236
## p_undergrad p_undergrad 1.436907
## books       books     1.415833
## top25perc   top25perc  1.361722
## enroll      enroll    1.037453
```

The most important predictors are expend and room_board.

Report the test error.

```
pred.gbm <- predict(gbm1, newdata = test_df)
te_gbm = RMSE(pred.gbm, test_df$outstate)
te_gbm
```

```
## [1] 1620.551
```

The test error is 1620.5511732.

Question 2

```
# Import and clean the data
data(OJ)
oj_df=
  OJ %>%
  na.omit() %>%
  janitor::clean_names()

# Partition data into training/test sets
set.seed(2)
oj_train = createDataPartition(y = oj_df$purchase,
                                p = 700/1070,
                                list = FALSE)

train_oj = oj_df[oj_train,]
test_oj = oj_df[-oj_train,]
```

Part a


Build a classification tree using the training data

```
set.seed(2)
tree2 <- rpart(formula = purchase ~ . ,
               data = train_oj,
               control = rpart.control(cp = 0))

cpTable <- printcp(tree2)
```

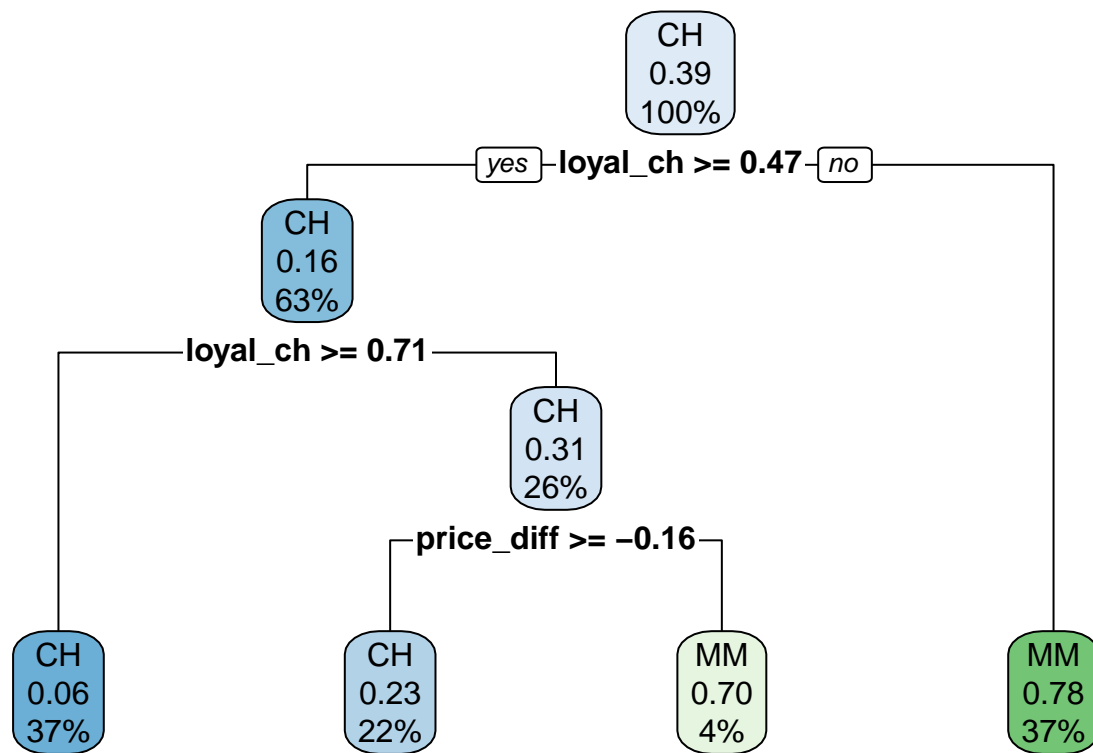
```
##
## Classification tree:
## rpart(formula = purchase ~ ., data = train_oj, control = rpart.control(cp = 0))
##
## Variables actually used in tree construction:
## [1] list_price_diff loyal_ch      price_diff      sale_price_ch
## [5] sale_price_mm   store      store_id      weekof_purchase
##
## Root node error: 273/701 = 0.38944
##
## n= 701
##
##      CP nsplit rel error  xerror    xstd
## 1 0.5201465      0  1.00000 1.00000 0.047291
## 2 0.0219780      1  0.47985 0.49451 0.038243
## 3 0.0097680      3  0.43590 0.48352 0.037916
## 4 0.0073260      6  0.40659 0.50183 0.038457
## 5 0.0045788      8  0.39194 0.48352 0.037916
## 6 0.0036630     12  0.37363 0.49451 0.038243
## 7 0.0012210     18  0.34799 0.49451 0.038243
## 8 0.0000000     21  0.34432 0.50916 0.038668
```

```
plotcp(tree2)
```



Use cross-validation to determine the tree size and create a plot of the final tree.

```
# minimum cross-validation error
minErr <- which.min(cpTable[,4])
tree3 <- prune(tree2, cp = cpTable[minErr,1])
rpart.plot(tree3)
```



```
summary(tree3)
```

```
## Call:
## rpart(formula = purchase ~ ., data = train_oj, control = rpart.control(cp = 0))
##   n= 701
##
##           CP nsplit rel error   xerror   xstd
## 1 0.52014652      0 1.0000000 1.0000000 0.04729133
## 2 0.02197802      1 0.4798535 0.4945055 0.03824313
## 3 0.00976801      3 0.4358974 0.4835165 0.03791591
##
## Variable importance
##      loyal_ch      price_diff  sale_price_mm  pct_disc_mm      disc_mm
##           77           7           4           3           3
## weekof_purchase      price_ch  special_mm      price_mm      store_id
##           2           1           1           1           1
##
## Node number 1: 701 observations,   complexity param=0.5201465
##   predicted class=CH   expected loss=0.3894437   P(node) =1
##   class counts:   428   273
##   probabilities: 0.611 0.389
##   left son=2 (443 obs) right son=3 (258 obs)
##   Primary splits:
##     loyal_ch < 0.469289 to the right, improve=121.50000, (0 missing)
##     store_id < 3.5      to the right, improve= 32.47240, (0 missing)
```

```

##      price_diff < 0.015      to the right, improve= 21.38773, (0 missing)
##      store7      splits as RL,      improve= 19.81980, (0 missing)
##      store      < 0.5      to the left, improve= 19.81980, (0 missing)
## Surrogate splits:
##      disc_mm      < 0.57      to the left, agree=0.641, adj=0.023, (0 split)
##      pct_disc_mm  < 0.264375 to the left, agree=0.641, adj=0.023, (0 split)
##      sale_price_mm < 1.385      to the right, agree=0.638, adj=0.016, (0 split)
##      price_diff    < -0.575    to the right, agree=0.638, adj=0.016, (0 split)
##      sale_price_ch < 2.025      to the left, agree=0.633, adj=0.004, (0 split)
##
## Node number 2: 443 observations,      complexity param=0.02197802
## predicted class=CH expected loss=0.1647856 P(node) =0.6319544
## class counts: 370 73
## probabilities: 0.835 0.165
## left son=4 (259 obs) right son=5 (184 obs)
## Primary splits:
##      loyal_ch      < 0.705699 to the right, improve=13.233360, (0 missing)
##      price_diff     < -0.39      to the right, improve=11.999940, (0 missing)
##      sale_price_mm  < 2.04      to the right, improve= 7.131761, (0 missing)
##      special_mm     < 0.5      to the left, improve= 5.427963, (0 missing)
##      list_price_diff < 0.235      to the right, improve= 5.218477, (0 missing)
## Surrogate splits:
##      price_ch      < 1.775      to the right, agree=0.639, adj=0.130, (0 split)
##      weekof_purchase < 237.5      to the right, agree=0.634, adj=0.120, (0 split)
##      price_mm      < 2.04      to the right, agree=0.630, adj=0.109, (0 split)
##      store_id      < 2.5      to the right, agree=0.614, adj=0.071, (0 split)
##      sale_price_mm  < 2.04      to the right, agree=0.605, adj=0.049, (0 split)
##
## Node number 3: 258 observations
## predicted class=MM expected loss=0.2248062 P(node) =0.3680456
## class counts: 58 200
## probabilities: 0.225 0.775
##
## Node number 4: 259 observations
## predicted class=CH expected loss=0.06177606 P(node) =0.3694722
## class counts: 243 16
## probabilities: 0.938 0.062
##
## Node number 5: 184 observations,      complexity param=0.02197802
## predicted class=CH expected loss=0.3097826 P(node) =0.2624822
## class counts: 127 57
## probabilities: 0.690 0.310
## left son=10 (154 obs) right son=11 (30 obs)
## Primary splits:
##      price_diff     < -0.165      to the right, improve=10.915950, (0 missing)
##      list_price_diff < 0.235      to the right, improve= 8.137578, (0 missing)
##      store_id      < 5.5      to the right, improve= 5.270274, (0 missing)
##      store7      splits as RL,      improve= 5.270274, (0 missing)
##      store      < 0.5      to the left, improve= 5.270274, (0 missing)
## Surrogate splits:
##      sale_price_mm  < 1.585      to the right, agree=0.891, adj=0.333, (0 split)
##      pct_disc_mm    < 0.187437 to the left, agree=0.886, adj=0.300, (0 split)
##      disc_mm      < 0.57      to the left, agree=0.880, adj=0.267, (0 split)
##      weekof_purchase < 274.5      to the left, agree=0.875, adj=0.233, (0 split)


```

```
##      special_mm      < 0.5      to the left,  agree=0.859, adj=0.133, (0 split)
##
## Node number 10: 154 observations
##   predicted class=CH   expected loss=0.2337662   P(node) =0.2196862
##   class counts:    118    36
##   probabilities: 0.766 0.234
##
## Node number 11: 30 observations
##   predicted class=MM   expected loss=0.3   P(node) =0.04279601
##   class counts:      9    21
##   probabilities: 0.300 0.700
```

When the tree size is 3 corresponds to the lowest cross-validation error.

Using the 1 SE rule

```
tree4 <- prune(tree2, cp = cpTable[cpTable[,4]<cpTable[minErr,4]+cpTable[minErr,5],1][1])
rpart.plot(tree4)
```



The tree is smaller when the tree size obtained using the 1 SE rule.

Part b


Perform boosting on the training data.

```
ctrl2 <- trainControl(method = "cv",
                      classProbs = TRUE,
                      summaryFunction = twoClassSummary)

gbmA.grid <- expand.grid(n.trees = c(2000,3000,4000,5000),
                       interaction.depth = 1:6,
                       shrinkage = c(0.0005,0.001,0.002),
                       n.minobsinnode = 1)


set.seed(2)
gbmA.fit <- train(purchase ~ . ,
                 data = train_oj,
                 tuneGrid = gbmA.grid,
                 trControl = ctrl2,
                 method = "gbm",
                 distribution = "adaboost",
                 metric = "ROC",
                 verbose = FALSE)

ggplot(gbmA.fit, highlight = TRUE)
```



Report the variable importance.

```
summary(gbmA.fit$finalModel, las = 2, cBars = 16, cex.names = 0.6)
```



```
##           var    rel.inf
## loyal_ch      loyal_ch 71.5077135
## price_diff    price_diff 12.5439501
## list_price_diff list_price_diff 2.4675725
## sale_price_mm  sale_price_mm 2.4288086
## store_id      store_id 2.3498719
## weekof_purchase weekof_purchase 2.2994427
## store         store 1.4990614
## disc_ch       disc_ch 1.1872920
## price_mm      price_mm 0.6534050
## pct_disc_ch   pct_disc_ch 0.6014305
## sale_price_ch sale_price_ch 0.5990112
## price_ch      price_ch 0.5322951
## special_mm    special_mm 0.3947971
## store7Yes     store7Yes 0.3494613
## special_ch    special_ch 0.2865049
## disc_mm       disc_mm 0.1970333
## pct_disc_mm   pct_disc_mm 0.1023489
```

The most important variables are loyal_ch and price_diff.

Test error

```
gbmA.pred <- predict(gbmA.fit, newdata = test_oj)
test.error = mean(gbmA.pred != test_oj$purchase)
test_error_rate = test.error*100
test_error_rate
```

```
## [1] 16.53117
```

The test error rate is 16.5311653%.